

Basic Tools for NLP

Rishu Kumar, Iuliia Zaitova, Chris Hyek

Saarland University

20.07.2023

Outline: cluster computing

- ▶ Computer clusters at coli
- ▶ File system
- ▶ CPU / GPU / CUDA
- ▶ Requirements of your job
- ▶ Running jobs at the coli cluster

Coli account

Do you have it?

If not, *please* follow the link:

<https://wiki.coli.uni-saarland.de/public/AccountApplication>



Coli cluster

- ▶ Login to the cluster:

```
ssh your_coli_username@login.coli.uni-saarland.de
```

- ▶ See the output of the "rwho" and "ruptime" commands


```
uptime / load average on nodes (output of 'ruptime'):  
falken-1    down   12+00:04  
falken-2    up     1+23:52,    0 users,  load  0.05,  0.02,  0.00  
falken-4    up     8+00:55,    0 users,  load  0.04,  0.01,  0.05  
jones-1     up     1+23:52,    0 users,  load  1.00,  1.00,  1.01  
jones-2     up    158+20:50,    0 users,  load  2.31,  2.29,  2.23  
jones-3     up     64+19:49,    0 users,  load  2.00,  2.04,  2.09  
jones-4     up     64+02:41,    2 users,  load  9.00,  7.69,  7.58  
jones-5     up     11+21:16,    0 users,  load  1.07,  0.66,  0.28  
jones-6     up      8+20:07,    0 users,  load  0.00,  0.00,  0.00  
tony-1      up    247+01:57,    3 users,  load  4.87,  3.42,  2.23  
tony-2      up    146+22:14,    1 user,   load  1.49,  1.51,  1.68  
tony-3      up     76+23:16,    1 user,   load 61.13, 45.85, 42.92  
tony-4      up     64+21:57,    2 users,  load  8.26,  8.34,  8.44
```

What do these mean?

- ▶ **falken-1 to falken-4**: Large memory compute servers Dell R815 featuring 4 AMD Opteron 6380 16-core processors (64 core in total) at 1TB RAM, each.
- ▶ **fjones-1 to jones-6**: GPU compute servers Dell T630 featuring 4 Intel Xeon E5-2687W v3 10-core processors (40 cores in total) at 256GB RAM and 4 MSI nvidia Titan-X (2015) GPU cards (12GB), each.
- ▶ **ftony-1 (aka volta-1) and tony-2 (aka volta-2)**: GPU compute servers Supermicro featuring 2 Intel Xeon 6128 6-core processors (12 cores in total, 24 cores with hyperthreading) at 768GB RAM and 8 nvidia Tesla V100 (2018) GPU cards (32GB), each.
- ▶ **ftony-3 (aka ampere-1)**: GPU compute servers Supermicro featuring 2 AMD EPYC 7713 64-core processors (128 cores in total, 256 cores with hyperthreading) at 512GB RAM and 8 nvidia Tesla A100 (2022) GPU cards (80GB).
- ▶ **ftony-4 (aka ampere-2)**: GPU compute servers Supermicro featuring 2 AMD EPYC 7662 64-core processors (128 cores in total, 256 cores with hyperthreading) at 512GB RAM and 8 nvidia Tesla A100 (2022) GPU cards (40GB), each.

CPU (Central Processing Unit)

- ▶ General-purpose processor designed to handle sequential tasks.
- ▶ Fewer but more powerful cores with complex instruction pipelines.




CPU
Multiple Cores



GPU (Graphics Processing Unit)

- ▶ Specialized for parallel processing of graphics and computational tasks.
- ▶ Thousands of cores designed to handle multiple threads simultaneously.



GPU
Thousands of Cores

CPU (Central Processing Unit)

- ▶ General-purpose processor designed to handle sequential tasks.
- ▶ Fewer but more powerful cores with complex instruction pipelines.
- ▶ **Used in general computing tasks, operating systems, and software applications.**

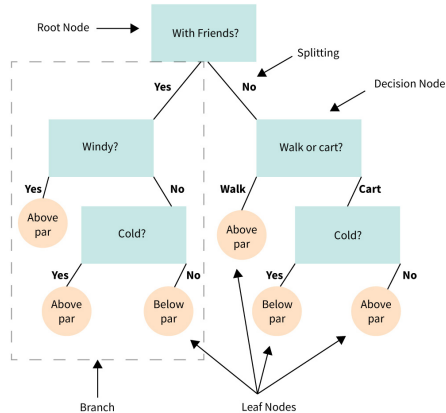
GPU (Graphics Processing Unit)

- ▶ Specialized for parallel processing of graphics and computational tasks.
- ▶ Thousands of cores designed to handle multiple threads simultaneously.
- ▶ **Optimized for data parallelism.**
- ▶ Often used in ML and AI (and gaming) — efficiently handles large matrices, vector operations.

3D graphics rendering?




Programs that contain a lot of conditional statements? (if, switch, while and such)



Applying filter to a photo?



Recursive algorithms?



When doing machine learning

1. Run serial workload on CPU
2. Offload parallel computation to GPUs

Instructions by Noon Pokaratsiri Goldstein



`iuliazaitova.github.io/basic-tools-nlp-2023`

Download the instructions and follow them

`pip install torch, numpy`

`run numpy_matrices.py` on the server

`run cpu_task.py` on the server

experiment with `nvidia-smi`

`export CUDA_VISIBLE_DEVICES=0,1`